

Introduction to Gaudi and Athena

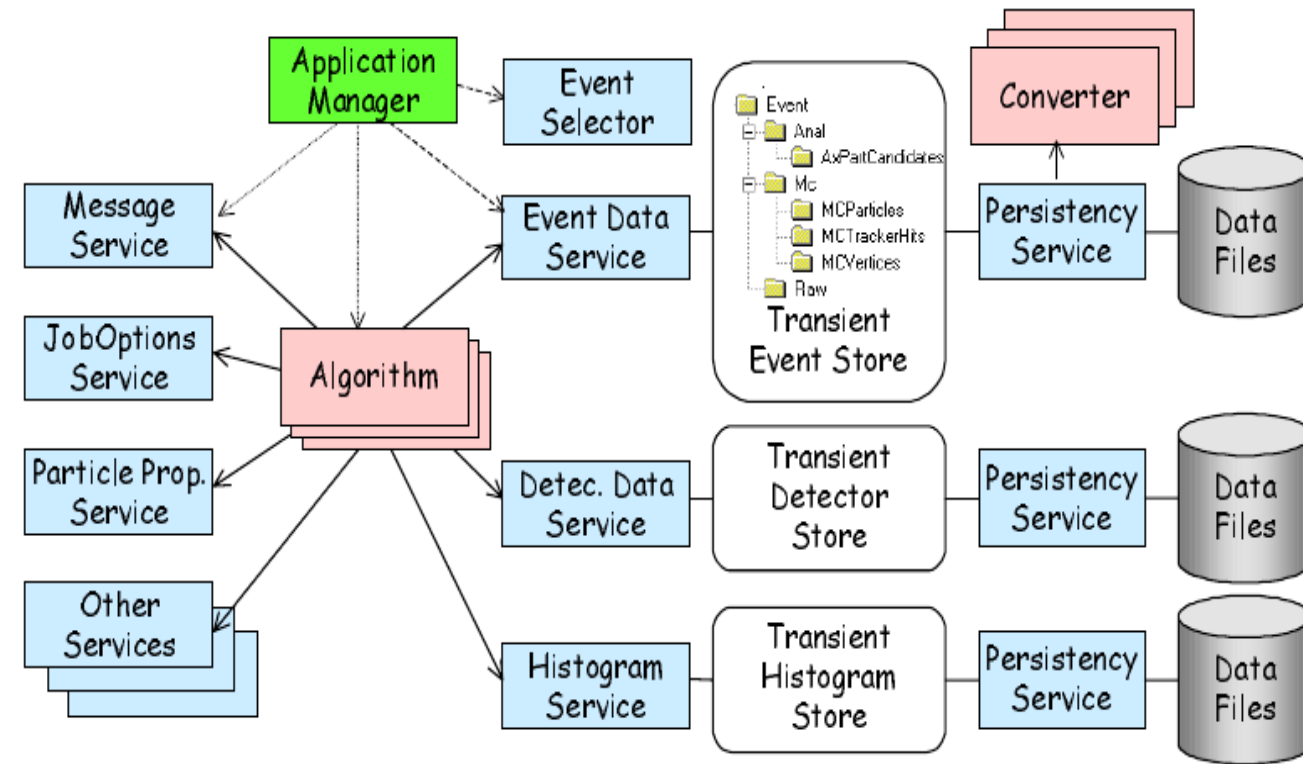
Vakho Tsulaia (LBL)

*Based on the presentations given by **Charles Leggett (LBL)** and **Paolo Calafiura (LBL)** in earlier editions of the ATLAS Software Development Tutorial*

ATLAS Software Development Tutorial
CERN, September 23-27, 2019

Building blocks of Gaudi

- **Algorithm**
 - Main building block of the **Event Loop**
 - Called once per event
- **AlgTool**
 - A plugin that helps an Algorithm perform some action
- **Service**
 - A plugin providing a common service to multiple components
 - **Examples:** Transient Data Store, Logging Service, Random Number Service



Interfaces, Plug-ins, Factories, etc.

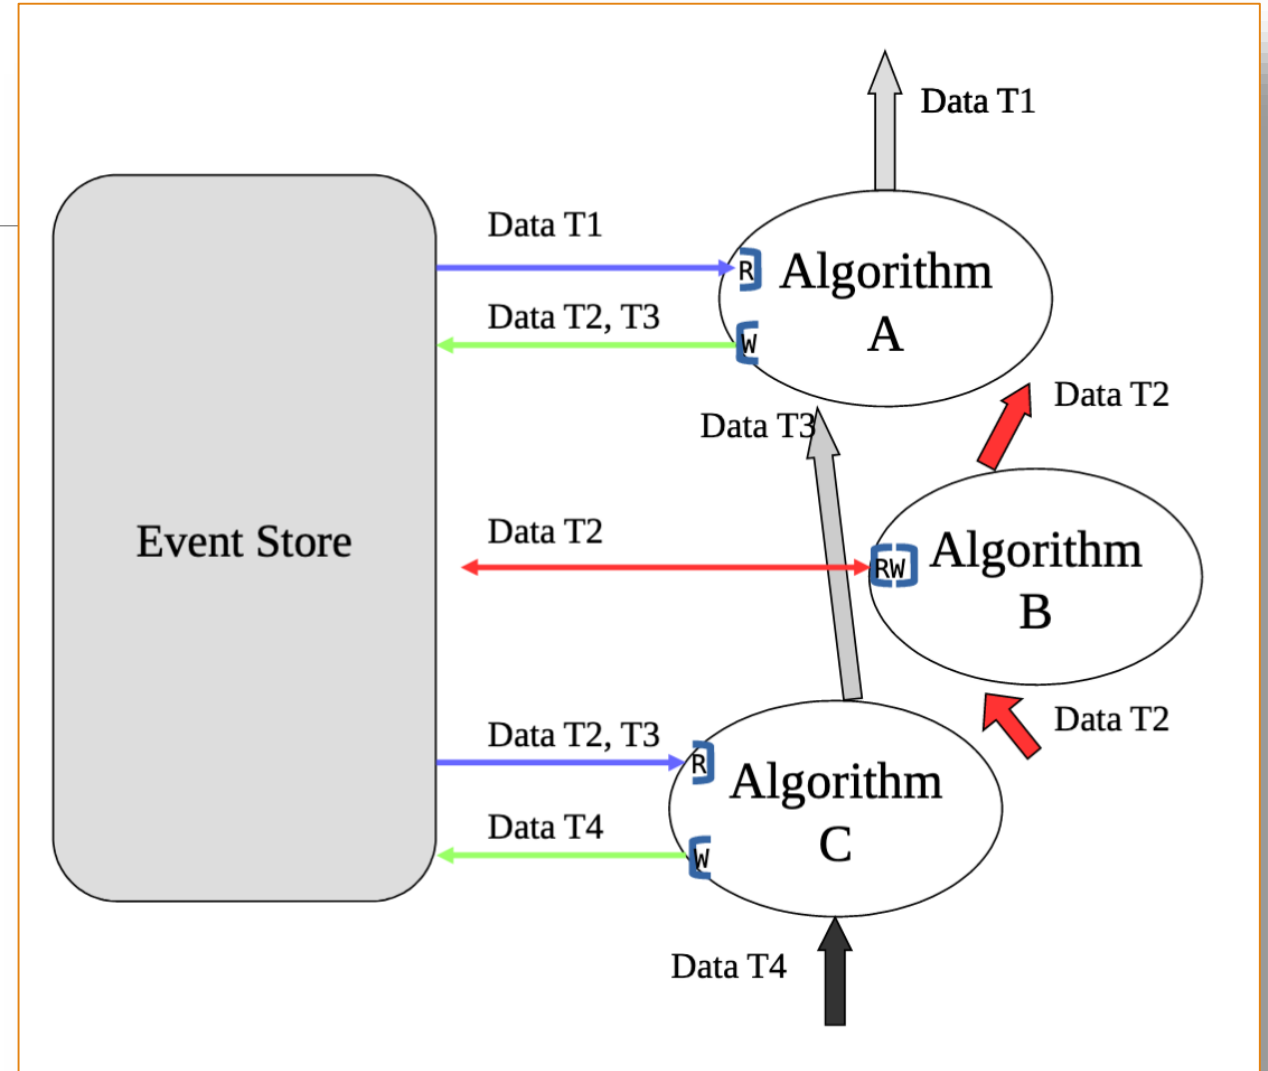
- Gaudi component model originally inspired by Microsoft COM
- Components implement **an interface** and use other components through **an interface**
- Components get packaged into Shared Object Libraries (DSO) and declared in a special component manifest files
- A dedicated Gaudi Service (**PluginSvc**) uses manifest to locate which DSO contains requested component dictionary, dl-opens it and creates an instance using a factory method

Job execution stages (simplified)

- **Configuration**
 - Parsing of configuration scripts (**Job Options**)
 - *More on configuration later this week*
- **Initialization**
 - `::initialize()` methods of components (Services, Algorithms) get called
- **Event processing loop**
 - `::execute()` methods of Algorithms get called
- **Finalization**
 - `::finalize()` methods of components (Services, Algorithms) get called

Event processing in serial Athena

- Algorithms run in sequences defined at the configuration time
- Algorithms consume (read) and produce (write) data objects shared via an Event Store
- Objects in the Event Store define data dependencies between algorithms
- Algorithms don't call each others methods



Status codes

- `StatusCode` is a standard return code used by Athena / Gaudi components
- Values are `StatusCode::SUCCESS` and `StatusCode::FAILURE`
- **Must be checked** on return of function. Unchecked `StatusCode` results in termination of the Athena job

```
StatusCode sc;  
sc = SomeFunction(par);  
If ( sc.isFailure() ) {  
    ATH_MSG_FATAL("SomeFunction failed for par : " << par);  
    return sc;  
}
```

Check StatusCode

```
ATH_CHECK(OtherFunction(par));
```

Recommended way of
checking StatusCode

Data access in Athena

- **StoreGateSvc** – **Transient Object Store** implementation in Athena
 - Manages data objects in memory (owns them)
 - Provides interface for **recording** and **retrieving** objects to/from the object store
 - Type-centric Naming Service. Each object has a **unique key**
- Several instances of Store Gate in typical Athena job
- Event Data:
 - **Event Store**
 - Contains data corresponding to one event. Gets cleared after every event
- Non-Event Data
 - **Detector Store** (immutable data, e.g. detector geometry)
 - **Metadata Store** (in-file metadata)
 - **Condition Store** (time-varying calibration data retrieved from Condition Database)

Data Handles

- Athena components access data in the Event Store via **Smart Handles**
- Two types of handles are defined for general use: **ReadHandle** and **WriteHandle**
- The handles are **templated** on the type of object being referenced
- They hold the corresponding **StoreGate** key
- **They act like pointers to the referenced data**, caching the pointer to the object in the Event Store

Usage of Handles

- Declare corresponding **Handle Keys** as private data members of your Algorithm/AlgTool

```
class MyAlg : public AthAlgorithm {  
...  
private:  
    SG::ReadHandleKey<MyObj1> m_myObj1Key {this, "MyObj1", "DefaultKey1", "MyObj1 Handle Key"};  
    SG::WriteHandleKey<MyObj2> m_myObj2Key {this, "MyObj2", "DefaultKey2", "MyObj2 Handle Key"};
```

- This automatically sets up handle key properties of the Algorithm, which can be set in job options

Default value

Property name

```
topseq += MyAlg (... , myObj1 = 'otherKey1')  
Topseq.MyAlg.myObj2 = 'otherKey2'
```

Usage of Handles (contd.)

- Initialize your handle keys from the algorithm's `initialize()` method

```
StatusCode MyAlg::initialize() {  
...  
    ATH_CHECK( m_myObj1Key.initialize() );  
    ATH_CHECK( m_myObj2Key.initialize() );  
}
```

- To use the handles in `execute()` of your algorithm, create handle objects as local variables, initializing them from the keys

```
StatusCode MyAlg::execute() {  
...  
    SG::ReadHandle<MyObj1> myObj1 (m_myObj1Key);  
    SG::WriteHandle<MyObj2> myObj2 (m_myObj2Key);  
}
```

Usage of Handles (contd.)

- To record the object to the store:

```
StatusCode MyAlg::execute() {  
...  
    SG::WriteHandle<MyObj2> myObj2 (m_myObj2Key);  
    ATH_CHECK( myObj2.record (std::make_unique<MyObj2>()) );  
}
```

- For in-depth documentation of the data access in Athena, see the [“Event data access in AthenaMT”](#) TWiki page

Accessing components

- Algorithms/AlgTools access Services and AlgTools via **Handles**
- Declare handles as data members of your Algorithm/AlgTool (ToolHandle **must** be private!)

```
class MyAlg : public AthAlgorithm {  
...  
private:  
    ToolHandle<IHiveTool>      m_hiveTool {this, "PrivateTool", "PriHiveTool", "private tool"};  
    ServiceHandle<IGeoModelSvc> m_geoModelSvc {this, "GeoModelSvc", "GeoModelSvc"};
```

- Retrieve them at initialize ...

```
StatusCode MyAlg::initialize() {  
...  
    ATH_CHECK(m_hiveTool.retrieve());  
    ATH_CHECK(m_geoModelSvc.retrieve());
```

- ... and then treat them like pointers

```
m_hiveTool->methodA();  
...  
m_geoModelSvc->methodB();
```

Configuring a component

- In order to be configurable via the Configuration layer, an Algorithm/AlgTool must **declare one or several properties**

```
class MyAlg : public AthAlgorithm {  
...  
private:  
    Gaudi::Property<bool> m_flag {this, "Flag", true, "some description"};  
    Gaudi::Property<float> m_float {this, "Float", 3.14, "looks like PI"};
```

Property name

Default value

Doc string

- You can set the property values in **job options** script

```
topseq += MyAlg (... , Flag = False)  
Topseq.MyAlg.Float = 3.1416
```

Const-correctness of ToolHandles

- In general, you should think of ToolHandles as objects, not pointer to objects, even though you will access the methods of the Tool with pointer-like semantics.
- ToolHandles are preserving constness: **it is not possible to call a non-const method of a Tool from a const ToolHandle**
- The following will generate a compile time error
 - If a class has a ToolHandle as a data member, calling a non-const method of the Tool from a const function of the class
 - Iterating over a ToolHandleArray with a const_iterator, and calling non-const methods
 - Extracting a T* pointer from a const ToolHandle<T>

Run Athena

- Once you have your job options script ready, the next step is to run the job:

```
$ athena.py [OPTIONS] jobOptions.py
```

... before doing that, though, first you need to setup runtime 😊 (see next slides) ...

- This is **not** how we run Athena in production
- For the production jobs we use special wrapper scripts called **Job Transforms**
 - Provide interface to the production system components (e.g. **Pilot**)
 - One transform can launch either a single Athena job or a chain of jobs, each implementing different data processing stages (e.g. digitization, reconstruction, trigger simulation)
 - **Not covered in this tutorial**

Releases

- [Athena repository](#) in GitLab has number of branches
 - master – main development branch, now leading to 22.0.X series of releases
 - 21.0, 21.1, 21.2, 21.3 – several production release branches
- For several branches (e.g. master, 21.X) we automatically build nightly releases
 - Each nightly build has corresponding Git tag
Example: `nightly/master/2019-09-21T2133`
 - Nightly builds are installed on CVMFS and kept around for ~1 month
- Status of nightly builds can be monitored [here](#)
- Once there is a need for that, we build **numbered releases** of a particular branch
Examples: `21.0.100`, `22.0.4`
 - Each of these releases has a corresponding tag in GitLab
 - Unlike nightly builds, the numbered releases are there to stay “forever”

Runtime

- Once you start a new session, say, on 1xp1us, first thing you usually do is to run this command

```
$ setupATLAS
```

- After that it is recommended to pick up “right” version of git

```
$ 1setup git
```

- Now you are all set to setup runtime against some recent nightly release

```
$ asetup Athena, master, r21
```

Nightly


Branch

Project

- ... or some numbered release

```
$ asetup Athena, 22.0.4
```


Athena MT



Gaudi MT

Parallel programming models

Problem decomposition viewpoint

- 
- Task-parallel model
 - Leverages **intra**-event concurrency
 - Data-parallel model
 - Leverages **inter**-event concurrency

Task-based programming


An abstraction of hardware architecture

- Faster task startup and shutdown
 - tthread/ttask = 18 onLinux (*)
 - tthread/ttask = 100 onWindows (*)
- Better load balancing
- Favors structured work partitioning and task precedence management
- Matching parallelism to resources
- Matching tasks to resources

(*) According to Intel Threading Building Blocks Team
tthread – time to launch a thread

AthenaMT

- Multi-threaded successor of Athena, designed to make better use of existing and upcoming hardware, while minimizing the memory footprint
- Uses the same component model as serial Athena with the same states (`initialize`, `execute`, `finalize`)
- Data flow driven: `algorithms declare what data they consume and produce`
- Each `Algorithm::execute()` happens in its own thread from shared thread pool
- Event pipelining: `multiple event can be executed concurrently`
- Algorithm cloning: `multiple instances of the same algorithm can execute with different event contexts`



Event processing:

- each event is a different color
- Each shape is a different algorithm

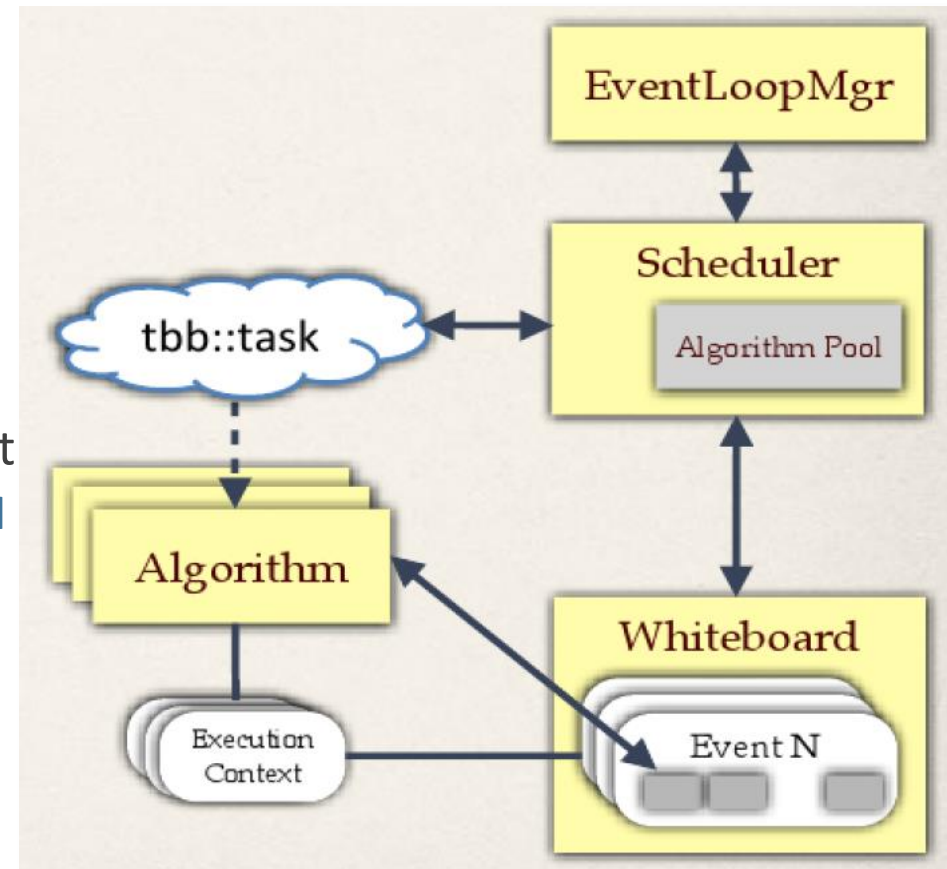
AthenaMT implementation



- AthenaMT uses **Intel Threaded Building Blocks (TBB)** for thread management
 - C++ templates and generic programming
 - Abstracts threading model by allowing operations to be treated as **tasks**
 - Supports **task stealing**: tasks on busy cores automatically reassigned to less busy ones
 - Library of algorithms and data structures that ease use of threads
 - `parallel_for`, `parallel_reduce`, `parallel_scan`, `parallel_do`
 - `concurrent_hash_map`, `concurrent_vector`, `concurrent_queue`
- While the TBB layer is hidden from Athena users, they can still use it for finer grained parallelism within Algorithms
 - In general not recommended

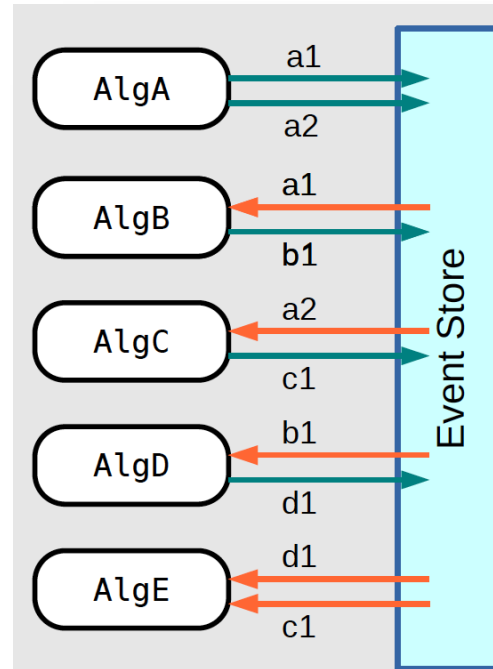
Athena MT operation

- Configuration, Initialization and Finalization are performed serially in the “master” thread
 - Only `Algorithm::execute()` is concurrent
- Algorithms are only scheduled when their input data becomes available
- Several instances of the same Algorithm can coexist
 - **Cloning:** create new instance if can be scheduled, and all other instances are busy
- Multiple events can be executed concurrently
 - Separate instance of Data Store per concurrent event




Scheduler

- Framework component that executes an Algorithm in an unused thread at the appropriate time
- Number of ways to decide which Algorithm to execute next
 - Precedence rules are created from a combination of **Data Flow** and **Control Flow**
 - Algorithms can be scheduled from any event in the queue
- Additional information can be provided to the Scheduler to help it determine most efficient way to get to the end

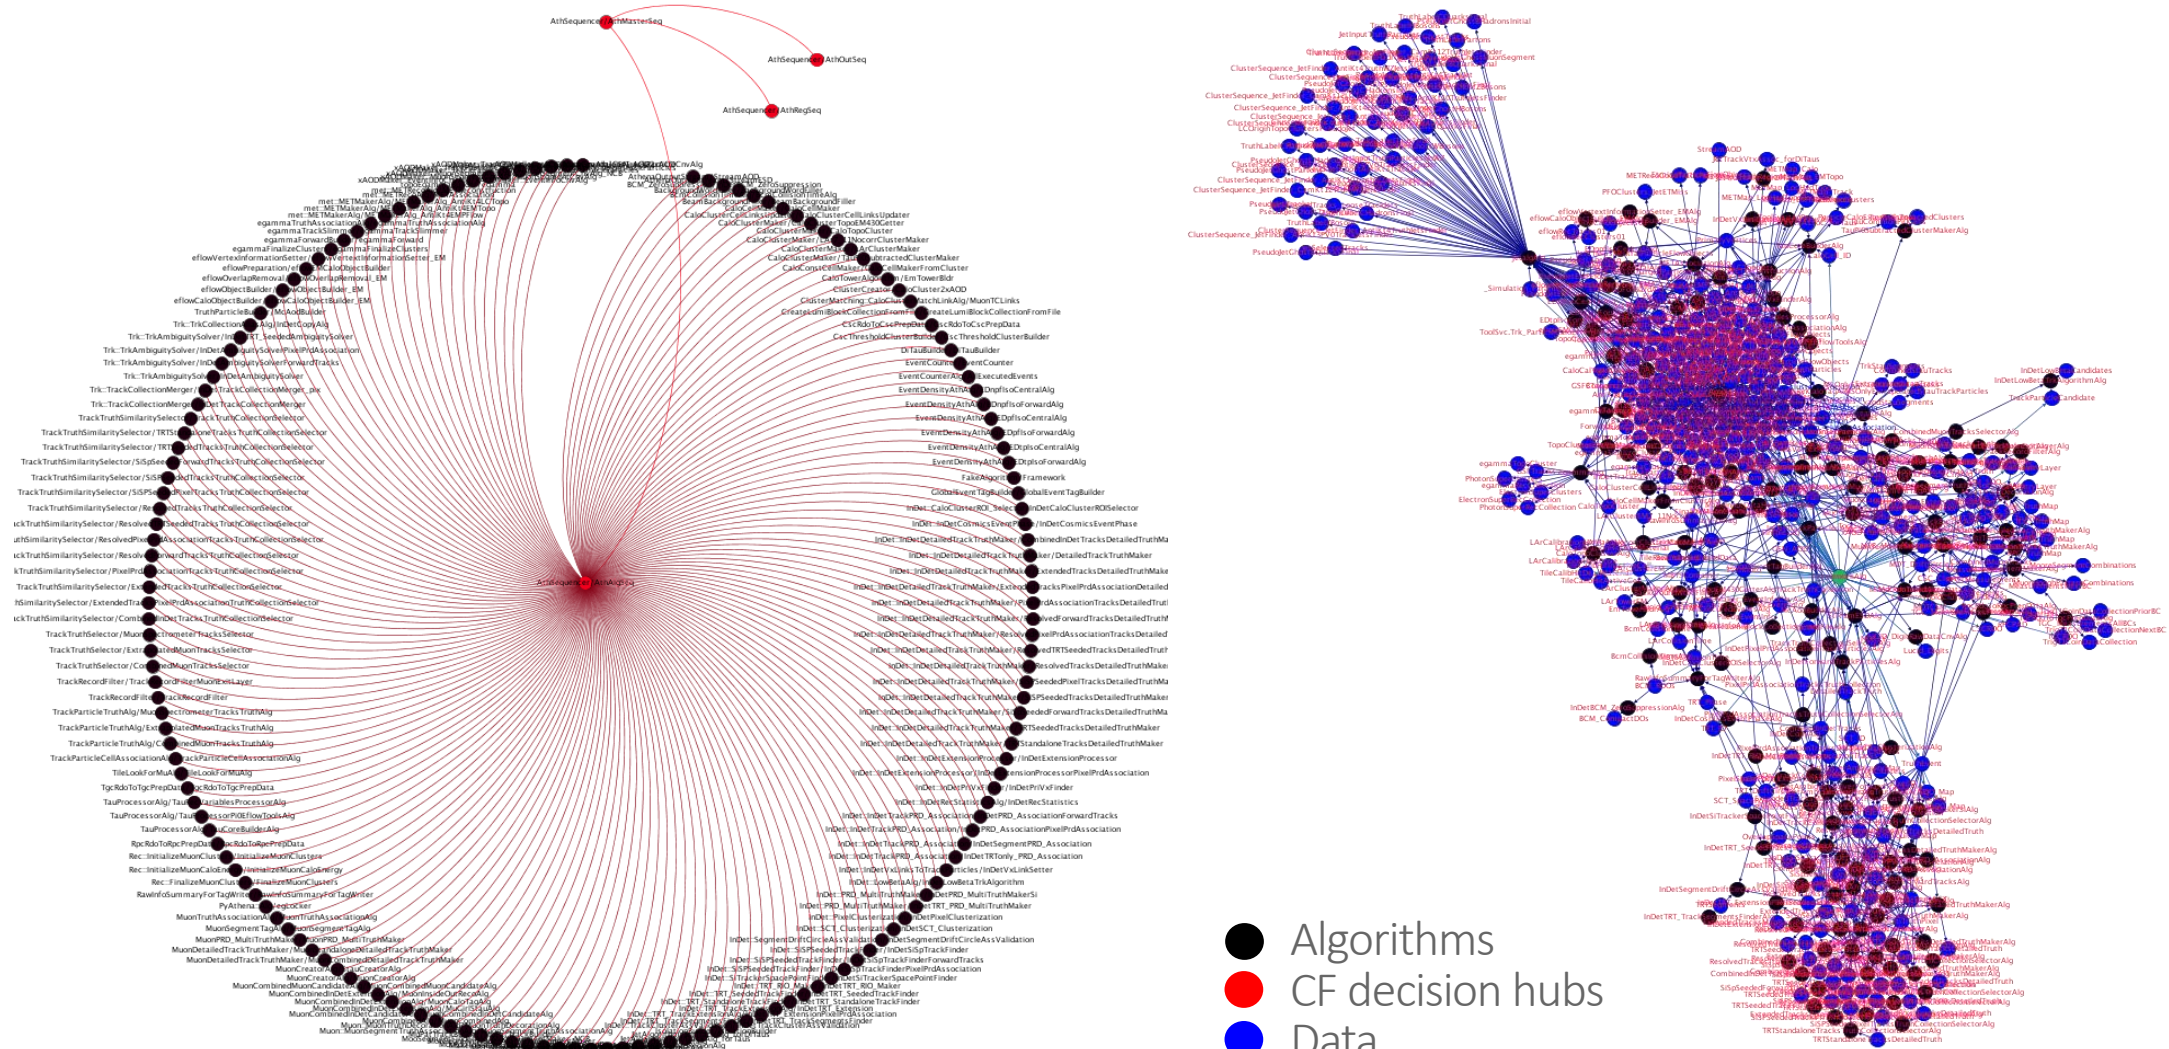


Serial



Concurrent

ATLAS MC data reconstruction precedence rules:
CF (left) and DF (right)



EventContext

- An object which provides information about the currently executed event
- Is set by the Event Loop Manager / Scheduler, and pushed into the Algorithm

```
#include "GaudiKernel/EventContext.h"
```

```
class EventContext {
```

```
public:
```

```
size_t evt() const;
```

```
size_t slot() const;
```

```
bool valid() const;
```

```
const EventIDBase& eventID() const;
```

```
const T* getExtension() const;
```

```
}
```

Number of events executed in job

Index of concurrent slot

Context is valid only within `Algorithm::execute`

Event#, Run#, LB#, Timestamp

Extendable content

Algorithm cloneability

- **Non-cloneable:** default
 - Only one instance exist in Algorithm Pool
 - Shared between all events. The Scheduler will **block** if algorithm needs to run on some event, but is in use in another event
 - Low memory use
- **Cloneable:**
 - Enabled explicitly by user via `MyAlg::isCloneable() { return true;}`
 - No clones created by default. Extra clones must be specifically requested by user
 - High memory use
- **Reentrant** (special base class `AthReentrantAlgorithm`):
 - **One instance**, shared by all threads
 - `execute()` is const and signature includes `EventContext`
 - Scheduler will **not block**, but use same instance for all threads
 - **Must be thread safe**

Component requirements

- **Algorithms:**
 - Must declare data dependencies (see later in this talk).
 - Data dependencies are union of Algorithm's and used AlgTools'
 - Communicate data to other Algorithms **only** via Whiteboard
 - **Must not cache** data between events
 - Should limit cloning to longer Algorithms, that are more likely to execute simultaneously
- **AlgTools** must be **stateless** and **private**
 - Must declare data dependencies
 - **Must not cache** data between events, unless in an **EventContext**-aware fashion
 - Can have event related data passed to them via accessors
 - Can access Whiteboard
- **Services** must be **thread safe** and **stateless**
 - May need to be made aware of the **EventContext**

Declaring data dependencies

- Algorithms (and AlgTools) must declare their Input and Output data dependencies to the framework
- Algorithms inherit the data dependencies of their sub-algorithms and Tools
 - Tools must be declared as **ToolHandles** by the Algorithms
- Data should be declared in the form of **VarHandleKeys**, using an initializer list syntax in the header of the Algorithm/AlgTool

```
class MyAlg : public AthAlgorithm {  
...  
private:  
    SG::ReadHandleKey<MyObj1> m_myObj1Key {this, "MyObj1", "DefaultKey1", "MyObj1 Handle Key"};  
    SG::WriteHandleKey<MyObj2> m_myObj2Key {this, "MyObj2", "DefaultKey2", "MyObj2 Handle Key"};
```

Examining job configuration

- You can dump Control Flow and Data Flow configuration of the job into log file

```
from AthenaCommon.AlgScheduler import AlgScheduler
AlgScheduler.ShowControlFlow( True )
AlgScheduler.ShowDataDependencies( True )
```

Examining job configuration (contd.)

- Data dependencies

```
AvalancheSchedulerSvc          0      INFO Data Dependencies for Algorithms:
  BeginIncFiringAlg
    none
  IncidentProcAlg1
    none
  xAODMaker::EventInfoCnvAlg
    o INPUT   ( 'EventInfo' , 'StoreGateSvc+McEventInfo' )
    o OUTPUT  ( 'SG::AuxElement' , 'StoreGateSvc+EventInfo' )
    o OUTPUT  ( 'SG::AuxVectorBase' , 'StoreGateSvc+PileupEventInfo' )
    o OUTPUT  ( 'xAOD::EventInfo' , 'StoreGateSvc+EventInfo' )
    o OUTPUT  ( 'xAOD::EventInfoContainer' , 'StoreGateSvc+PileupEventInfo' )
  SGInputLoader
    none
  HiveAlgA
    o INPUT   ( 'HiveDataObj' , 'StoreGateSvc+b2' )
    o INPUT   ( 'xAOD::EventInfo' , 'StoreGateSvc+EventInfo' )
    o OUTPUT  ( 'HiveDataObj' , 'StoreGateSvc+a1' )
    o OUTPUT  ( 'HiveDataObj' , 'StoreGateSvc+a2' )
  HiveAlgB
    o OUTPUT  ( 'HiveDataObj' , 'StoreGateSvc+b1' )
    o OUTPUT  ( 'HiveDataObj' , 'StoreGateSvc+b2' )
  HiveAlgC
    o INPUT   ( 'HiveDataObj' , 'StoreGateSvc+a1' )
    o OUTPUT  ( 'HiveDataObj' , 'StoreGateSvc+c1' )
    o OUTPUT  ( 'HiveDataObj' , 'StoreGateSvc+c2' )
```

Examining job configuration (contd.)

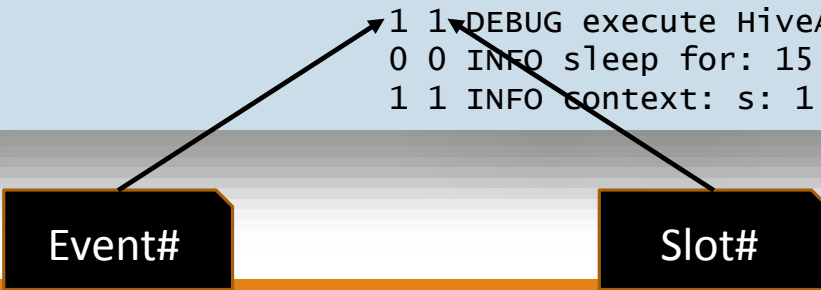
- Control Flow

```
===== Control Flow Configuration =====  
  
AthMasterSeq [Seq] [Sequential] [Prompt]  
  AthAlgEvtSeq [Seq] [Sequential]  
    AthBeginSeq [Seq] [Sequential] [Prompt]  
      BeginIncFiringAlg [Alg] [n= 0]  
      IncidentProcAlg1 [Alg] [n= 1] [unclonable]  
    AthAllAlgSeq [Seq] [Concurrent]  
      AthAlgSeq [Seq] [Concurrent] [PASS]  
        xAODMaker::EventInfoCnvAlg [Alg] [n= 0]  
        SGInputLoader [Alg] [n= 4]  
        HiveAlgA [Alg] [n= 4]  
        HiveAlgB [Alg] [n= 4]  
        HiveAlgC [Alg] [n= 4]  
        HiveAlgD [Alg] [n= 4]  
        HiveAlgE [Alg] [n= 4]  
        HiveAlgG [Alg] [n= 4]  
        HiveAlgF [Alg] [n= 4]  
        HiveAlgV [Alg] [n= 4]  
      AthCondSeq [Seq] [Concurrent]  
    AthEndSeq [Seq] [Sequential] [Prompt]  
      EndIncFiringAlg [Alg] [n= 0]  
      IncidentProcAlg2 [Alg] [n= 1] [unclonable]  
  AthOutSeq [Seq] [Concurrent]  
  AthRegSeq [Seq] [Concurrent]
```

Logging

- **MessageSvc** - a service used by Athena for logging purposes
- In AthenaMT, **MessageSvc** queues messages from each source
 - Atomic output, so log is readable
- By default you'll get slot# and event# (only if there is a valid **EventContext**)
 - **MessageSvc** supports several flags for output formatting (not covered in this talk)

```
AthenaHiveEventLoopMgr      3 3 INFO ===>>> start processing event #3, run #1 on slot 3, 0 events processed so far <<<===
AthenaHiveEventLoopMgr      INFO ===>>> done processing event #0, run #1 on slot 0, 1 events processed so far <<<===
AthenaHiveEventLoopMgr      4 4 INFO ===>>> start processing event #4, run #1 on slot 4, 1 events processed so far <<<===
HiveAlgA                    0 0 DEBUG execute HiveAlgA
HiveAlgB                    0 0 DEBUG execute HiveAlgB
HiveAlgB                    0 0 INFO context: s: 0 e: 0 for 0x17dc0000
HiveAlgB                    1 1 DEBUG execute HiveAlgB
HiveAlgB                    0 0 INFO sleep for: 15 ms
HiveAlgB                    1 1 INFO context: s: 1 e: 1 for 0x17dc4800
```



Event#

Slot#

Job Control

- Run an MT job like any other Athena job, and specify number of threads and/or number of concurrent events via Command-Line Interface
- Two knobs to tweak
 - `--threads=N` number of threads in Thread Pool
 - `--concurrent-events=N` number of concurrent events
- In general, easiest to just specify `--threads=N`, which will set both parameters to the same number
- Setting `--threads=1` will run **AthenaMT with 1 thread, not serial Athena**